

TRAIN THE ML MODEL ON IBM

Team ID	PNT2022TMID16353
Project Name	Car Resale value Prediction

TRAIN THE ML MODEL ON IBM

IBM Watson Studio

Search in your workspaces

Buy

?

1

SUGANTH V's Account

Dallas

SV

Deployments /

CAR

OverviewAssetsDeploymentsJobsManage

▼

Q

Search

↺

Name	Type	Status	Asset	Last modified	↓
Model_IBM	Online	✓ Deployed	resale_model	2 hours ago SUGANTH V (You)	⋮

Items per page: 201-1 of 1 items1 of 1 pages

```
import pandas as pd
import numpy as np
import matplotlib as plt
from sklearn.preprocessing import LabelEncoder
import pickle
print("IMPORTED REQUIRED LIBRARIES")
# df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0 , sep=',',
,encoding='Latin1',low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RKK',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
```

```

config=Config(signature_version='oauth'),
endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',encoding='Latin1',low_memory=False)
df.head()
# df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0 , sep=',',
.encoding='Latin1',low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-IL0017uhnUGwXyhG_Eort5gohoW6XJTNoT3RKK',
    ibm_auth_endpoint="https://iam.cloud.ibm.com/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv((io.BytesIO(body['Body'].read())) , header=0 , sep=',',encoding='Latin1',low_memory=False)
df.head()
print(df.seller.value_counts())
df[df.seller !='gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType !='Gesuch']
df=df.drop('offerType',axis=1)
print(df.shape)
df=df[(df.powerPS>50) & (df.powerPS<900)]
print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]
print(df.shape)
df.drop(['name','abtest','dateCrawled','nrOfPictures','lastSeen','postalCode','dateCreated'], axis='columns',inplace=True)
new_df=df.copy()
new_df=new_df.drop_duplicates(['price','vehicleType','yearOfRegistration','gearbox','powerPS','model','kilometer','monthOfRegistration','fuelType','notRepairedDamage'])
new_df.gearbox.replace(('manuell','automatic'),('manual','automatic'),inplace=True)
new_df.fuelType.replace(('benzin','andere','elektro'),('petrol','others','electric'),inplace=True)
new_df.vehicleType.replace(('kleinwagen','cabrio','kombi','andere'),('small car','convertible','combination','others'),inplace=True)

```

```

new_df.notRepairedDamage.replace(('ja','nein'),('Yes','No'),inplace=True)
new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

new_df['notRepairedDamage'].fillna(value='not-declared',inplace=True)
new_df['fuelType'].fillna(value='not-declared',inplace=True)
new_df['gearbox'].fillna(value='not-declared',inplace=True)
new_df['vehicleType'].fillna(value='not-declared',inplace=True)
new_df['model'].fillna(value='not-declared',inplace=True)
from ibm_watson_machine_learning import APIClient
wml_credentials={
    "url":"https://us-south.ml.cloud.ibm.com",
    "apikey":"hEAn_mcoP3u_-ZjagjeqlxDayqUiETpYVYWdR1OLKAby"
}
client =APIClient(wml_credentials)
def guide_from_space_name(client, space_name):
    space = client.spaces.get_details()
    # print(space)
    return(next(item for item in space['resources'] if item['entity']['name']==space_name)['metadata']['id'])
space_uid=guide_from_space_name(client,'CAR')
print("Space UID"+ space_uid)
client.set.default_space(space_uid)
client.software_specifications.list()
software_spec_uid = client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid
print(new_df)
labels=['gearbox','notRepairedDamage','model','brand','fuelType','vehicleType']

mapper={}
for i in labels:
    mapper[i]=LabelEncoder()
    mapper[i].fit(new_df[i])
    tr=mapper[i].transform(new_df[i])
    np.save(str('classes'+i+'.npy'),mapper[i].classes_)
    print(i,":",mapper[i])
    new_df.loc[:, i+ '_labels']=pd.Series(tr,index=new_df.index)

labeled = new_df[['price','yearOfRegistration','powerPS','kilometer','monthOfRegistration']+[x+" _labels" for x in
labels]]
print(labeled.columns)
Y=labeled.iloc[:,0].values
X=labeled.iloc[:,1:].values

Y=Y.reshape(-1,1)
from sklearn.model_selection import cross_val_score,train_test_split
X_train , X_test, Y_train , Y_test = train_test_split(X,Y,test_size=0.3,random_state=3)
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)

regressor.fit(X_train, np.ravel(Y_train,order='C'))
y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))
filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))

```

```

model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)
model_id
X_train[0]
regressor.predict([[2012.0, 179.0, '1500000', 12.0, 0, 0, 30, 1, 1, 4]])

```

IBM Watson Studio

Search in your workspaces

Buy

SUGANTH V's Account

Dallas

SV

Projects / CAR RESALE VALUE PREDICTION / Model Building

In [3]:

```

import pandas as pd
import numpy as np
import matplotlib as plt
from sklearn.preprocessing import LabelEncoder
import pickle
print("IMPORTED REQUIRED LIBRARIES")

```

IMPORTED REQUIRED LIBRARIES

In [4]:

```

# df = pd.read_csv("C:/Users/SUGARANJAN/Desktop/IBM/Data/autos.csv", header=0, sep=',', encoding='Latin1', low_memory=False)
# df.head()
import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3
import io
def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='DT151-1L0017uHnUGwXyhG_Eort5gohow6XJTNoT3RKK',
    ibm_auth_endpoint='https://iam.cloud.ibm.com/oidc/token',
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalueprediction-donotdelete-pr-yuhtmzidi0ka1p'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket,Key=object_key)
df = pd.read_csv(io.BytesIO(body['Body'].read()), header=0, sep=',', encoding='Latin1', low_memory=False)
df.head()

```

Out[4]:

	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand
0	24-03-2016 11.52	Golf_3_1.6	privat	Angebot	480.0	test	NaN	1993.0	manuell	0.0	golf	150000	0.0	benzin	volkswagen

Out[4]:

	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand
0	24-03-2016 11.52	Golf_3_1.6	privat	Angebot	480.0	test	NaN	1993.0	manuell	0.0	golf	150000	0.0	benzin	volkswagen
1	24-03-2016 10.58	A5_Sportback_2.7_Tdi	privat	Angebot	18300.0	test	coupe	2011.0	manuell	190.0	NaN	125000	5.0	diesel	audi
2	14-03-2016 12.52	Jeep_Grand_Cherokee_"Overland"	privat	Angebot	9800.0	test	suv	2004.0	automatik	163.0	grand	125000	8.0	diesel	jeep
3	17-03-2016 16.54	GOLF_4_1.4_3TURER	privat	Angebot	1500.0	test	kleinwagen	2001.0	manuell	75.0	golf	150000	6.0	benzin	volkswagen
4	31-03-2016 17.25	Skoda_Fabia_1.4_TDI_PD_Classic	privat	Angebot	3600.0	test	kleinwagen	2008.0	manuell	69.0	fabia	90000	7.0	diesel	skoda

In [5]:

```

print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'Gesuch']
df=df.drop('offerType',axis=1)

privat      371534
gewerblich    3
golf          1
Name: seller, dtype: int64
Angebot      371525
Gesuch        12
150000        1
Name: offerType, dtype: int64

```

In [6]: print(df.shape)

```

In [5]: print(df.seller.value_counts())
df[df.seller != 'gewerblich']
df=df.drop('seller',axis=1)

print(df.offerType.value_counts())
df[df.offerType != 'Gesuch']
df=df.drop('offerType',axis=1)

privat      371534
gewerblich      3
golf          1
Name: seller, dtype: int64
Angebot      371525
Gesuch       12
150000       1
Name: offerType, dtype: int64

In [6]: print(df.shape)
df=df[(df.powerPS>50) & (df.powerPS<900)]
print(df.shape)
df=df[(df.yearOfRegistration>=1950)&(df.yearOfRegistration<2022)]
print(df.shape)

(371539, 18)
(319717, 18)
(319649, 18)

In [7]: df.drop(['name', 'abtest', 'dateCrawled', 'nrOfPictures', 'lastSeen', 'postalCode', 'dateCreated'], axis='columns', inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price', 'vehicleType', 'yearOfRegistration', 'gearbox', 'powerPS', 'model', 'kilometer', 'monthOfRegistration', 'fuelType', 'notRepairedDamage'])

In [9]: new_df.gearbox.replace(('manuell', 'automatik'), ('manual', 'automatic'), inplace=True)
new_df.fuelType.replace(('benzin', 'andere', 'elektro'), ('petrol', 'others', 'electric'), inplace=True)
new_df.vehicleType.replace(('kleinwagen', 'cabrio', 'kombi', 'andere'), ('small car', 'convertible', 'combination', 'others'), inplace=True)
new_df.notRepairedDamage.replace(('ja', 'nein'), ('Yes', 'No'), inplace=True)

In [10]: new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

```

```

In [7]: df.drop(['name', 'abtest', 'dateCrawled', 'nrOfPictures', 'lastSeen', 'postalCode', 'dateCreated'], axis='columns', inplace=True)

In [8]: new_df=df.copy()
new_df=new_df.drop_duplicates(['price', 'vehicleType', 'yearOfRegistration', 'gearbox', 'powerPS', 'model', 'kilometer', 'monthOfRegistration', 'fuelType', 'notRepairedDamage'])

In [9]: new_df.gearbox.replace(('manuell', 'automatik'), ('manual', 'automatic'), inplace=True)
new_df.fuelType.replace(('benzin', 'andere', 'elektro'), ('petrol', 'others', 'electric'), inplace=True)
new_df.vehicleType.replace(('kleinwagen', 'cabrio', 'kombi', 'andere'), ('small car', 'convertible', 'combination', 'others'), inplace=True)
new_df.notRepairedDamage.replace(('ja', 'nein'), ('Yes', 'No'), inplace=True)

In [10]: new_df=new_df[(new_df.price>=100)&(new_df.price<=150000)]

new_df['notRepairedDamage'].fillna(value='not-declared', inplace=True)
new_df['fuelType'].fillna(value='not-declared', inplace=True)
new_df['gearbox'].fillna(value='not-declared', inplace=True)
new_df['vehicleType'].fillna(value='not-declared', inplace=True)
new_df['model'].fillna(value='not-declared', inplace=True)

In [11]: from ibm_watson_machine_learning import APIClient
wml_credentials={
    "url": "https://us-south.ml.cloud.ibm.com",
    "apikey": "hEAn_mcoP3u_-ZjagjeqlxDayQUIETpVYVNdR1OLKAbY"
}
client =APIClient(wml_credentials)

In [12]: def guide_from_space_name(client, space_name):
    space = client.spaces.get_details()
    # print(space)
    return(next(item for item in space['resources'] if item['entity']['name']==space_name)['metadata']['id'])

In [13]: space_uid=guide_from_space_name(client, 'CAR')
print("Space UID"+ space_uid)

Space UIDbe467bbb-03a2-40e7-bf5a-91836e346951

In [14]: client.set.default_space(space_uid)

Out[14]: 'SUCCESS'

```

In [15]: client.software_specifications.list()

NAME	ASSET_ID	TYPE
default_py3.6	0062b8c9-8b7d-44a0-a9b9-46c416adcdb9	base
kernel-spark3.2-scala2.12	020d69ce-7ac1-5e68-ac1a-31189867356a	base
pytorch-onnx_1.3-py3.7-edt	069ea134-3346-5748-b513-49120e15d288	base
scikit-learn_0.20-py3.6	09c5a1d0-9c1e-4473-a344-eb7b665ff687	base
spark-mllib_3.0-scala_2.12	09f4cff0-90a7-5899-b9ed-1ef348aebdee	base
pytorch-onnx_rt22.1-py3.9	0b848dd4-e681-5999-be41-b5f6fccc6471	base
ai-function_0.1-py3.6	0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda	base
shiny-r3.6	0e6e79df-875e-4f24-8ae9-62dcc2148306	base
tensorflow_2.4-py3.7-horovod	1092590a-307d-563d-9b62-4eb7d64b3f22	base
pytorch_1.1-py3.6	10ac12d6-6b30-4ccd-8392-3e922c096a92	base
tensorflow_1.15-py3.6-ddl	111e41b3-de2d-5422-a4d6-bf776828c4b7	base
autoai-kb_rt22.2-py3.10	125b6d9a-5b1f-5e8d-972a-b251688ccf40	base
runtime-22.1-py3.9	12b83a17-24d8-5082-900f-0ab31fbfd3cb	base
scikit-learn_0.22-py3.6	154010fa-5b3b-4ac1-82af-4d5ee5abbc85	base
default_r3.6	1b70aec3-ab34-4b87-8aa0-a4a3c829ea36	base
pytorch-onnx_1.3-py3.6	1bc6029a-cc97-56da-b8e0-39c3880dbbe7	base
kernel-spark3.3-r3.6	1c9e5454-f216-59dd-a20e-474a5cdf5988	base
pytorch-onnx_rt22.1-py3.9-edt	1d362186-7ad5-5b59-8b6c-9d0880bde37f	base
tensorflow_2.1-py3.6	1eb25b84-d6ed-5dde-b6a5-3fbd1665666	base
spark-mllib_3.2	20047f72-0a98-58c7-9ff5-a77b012eb8f5	base
tensorflow_2.4-py3.8-horovod	217c16f6-178f-56bf-824a-b19f20564c49	base
runtime-22.1-py3.9-cuda	26215f05-08c3-5a41-a1b0-da66306ce658	base
do_py3.8	295addb5-9ef9-547e-9bf4-92ae3563e720	base
autoai-ts_3.8-py3.8	2aa0c932-798f-5ae9-abd6-15e0c2402fb5	base
tensorflow_1.15-py3.6	2b73a275-7cbf-420b-a912-eae7f436e0bc	base
kernel-spark3.3-py3.9	2b7961e2-e3b1-5a8c-a491-482c8368839a	base
pytorch_1.2-py3.6	2c8ef57d-2687-4b7d-acce-01f94976dac1	base
spark-mllib_2.3	2e51f700-bca0-4b0d-88dc-5c6791338875	base
pytorch-onnx_1.1-py3.6-edt	32983cea-3f32-4400-8965-dde874a8d67e	base
spark-mllib_3.0-py37	36507ebe-8770-59ba-ab2a-eafe787600e9	base
spark-mllib_2.4	390d21f8-e58b-4fac-9c55-d7ceda621326	base
autoai-ts_rt22.2-py3.10	396b2e83-0953-5b86-9a55-7ce1628a406f	base
xgboost_0.82-py3.6	39e31acd-5f30-41dc-ae44-60233c80306e	base
pytorch-onnx_1.2-py3.6-edt	40589d0e-7019-4e28-8daa-fb03b6f4fe12	base
pytorch-onnx_rt22.2-py3.10	40e73f55-783a-5535-b3fa-0c8b94291431	base
default_r36py38	41c247d3-45f8-5a71-b065-8580278facf0	base

In [16]: software_spec_uid = client.software_specifications.get_uid_by_name("runtime-22.1-py3.9")
software_spec_uid

Out[16]: '12b83a17-24d8-5082-900f-0ab31fbfd3cb'

In [17]: print(new_df)

	price	vehicleType	yearOfRegistration	gearbox	powerPS	\
1	18300.0	coupe	2011.0	manual	190.0	
2	9800.0	suv	2004.0	automatic	163.0	
3	1500.0	small car	2001.0	manual	75.0	
4	3600.0	small car	2008.0	manual	69.0	
5	650.0	limousine	1995.0	manual	102.0	
...	
371531	3200.0	limousine	2004.0	manual	225.0	
371535	1199.0	convertible	2000.0	automatic	101.0	
371536	9200.0	bus	1996.0	manual	102.0	
371537	3400.0	combination	2002.0	manual	100.0	
371538	28990.0	limousine	2013.0	manual	320.0	
	model	kilometer	monthOfRegistration	fuelType	brand	\
1	not-declared	125000	5.0	diesel	audi	
2	grand	125000	8.0	diesel	jeep	
3	golf	150000	6.0	petrol	volkswagen	
4	fabia	90000	7.0	diesel	skoda	
5	3er	150000	10.0	petrol	bmw	
...	
371531	leon	150000	5.0	petrol	seat	
371535	fortwo	125000	3.0	petrol	smart	
371536	transporter	150000	3.0	diesel	volkswagen	
371537	golf	150000	6.0	diesel	volkswagen	
371538	m_reihe	50000	8.0	petrol	bmw	
	notRepairedDamage					
1	Yes					
2	not-declared					
3	No					
4	No					
5	Yes					

Projects / CAR RESALE VALUE PREDICTION / Model Building

```
[200000 rows x 11 columns]

In [18]: labels=['gearbox','notRepairedDamage','model','brand','fuelType','vehicleType']

mapper={}
for i in labels:
    mapper[i]=LabelEncoder()
    mapper[i].fit(new_df[i])
    tr=mapper[i].transform(new_df[i])
    np.save(str('classes'+i+'.npy'),mapper[i].classes_)
    print(i,":",mapper[i])
    new_df.loc[:, i+ '_labels']=pd.Series(tr,index=new_df.index)

labeled = new_df[['price','yearOfRegistration','powerPS','kilometer','monthOfRegistration']+['x_'+_labels for x in labels]]
print(labeled.columns)

gearbox : LabelEncoder()
notRepairedDamage : LabelEncoder()
model : LabelEncoder()
brand : LabelEncoder()
fuelType : LabelEncoder()
vehicleType : LabelEncoder()
Index(['price', 'yearOfRegistration', 'powerPS', 'kilometer',
       'monthOfRegistration', 'gearbox_labels', 'notRepairedDamage_labels',
       'model_labels', 'brand_labels', 'fuelType_labels',
       'vehicleType_labels'],
      dtype='object')

In [19]: Y=labeled.iloc[:,0].values
X=labeled.iloc[:,1:].values

Y=Y.reshape(-1,1)

In [20]: from sklearn.model_selection import cross_val_score,train_test_split
X_train , X_test, Y_train , Y_test = train_test_split(X,Y,test_size=0.3,random_state=3)

In [21]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)
```

Projects / CAR RESALE VALUE PREDICTION / Model Building

```
In [21]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)

regressor.fit(X_train, np.ravel(Y_train,order='C'))

Out[21]: RandomForestRegressor(max_depth=10, n_estimators=1000, random_state=34)

In [22]: y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))

0.8310350387286918

In [23]: filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))

In [24]: model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)

In [25]: model_id

Out[25]: 'cd8479e0-66e4-454e-aece-4824fe9d71bd'

In [26]: X_train[0:]

Out[26]: array([[2005.0, 179.0, '150000', ..., 1, 1, 4],
 [1997.0, 60.0, '150000', ..., 38, 7, 4],
 [2003.0, 170.0, '150000', ..., 2, 7, 1],
 ...,
 [2009.0, 174.0, '125000', ..., 25, 7, 7],
 [2000.0, 136.0, '150000', ..., 20, 7, 1],
 [2013.0, 170.0, '40000', ..., 1, 7, 8]], dtype=object)

In [27]: regressor.predict([[2012.0, 179.0, '150000', 12.0, 0, 0, 30, 1, 1, 4]])
```

```
regressor = RandomForestRegressor(n_estimators = 1000,max_depth = 10,random_state = 34)
regressor.fit(X_train, np.ravel(Y_train,order='C'))
```

Out[21]: RandomForestRegressor(max_depth=10, n_estimators=1000, random_state=34)

```
In [22]: y_pred = regressor.predict(X_test)
print(r2_score(Y_test,y_pred))
0.8310350387286918
```

```
In [23]: filename='resale_model.sav'
pickle.dump(regressor,open(filename,'wb'))
```

```
In [24]: model_details = client.repository.store_model(model=regressor,meta_props={
    client.repository.ModelMetaNames.NAME: "resale_model",
    client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid,
    client.repository.ModelMetaNames.TYPE: "scikit-learn_1.0"
})
model_id = client.repository.get_model_id(model_details)
```

In [25]: model_id

Out[25]: 'cd8479e0-66e4-454e-aece-4824fe9d71bd'

In [26]: X_train[0:]

Out[26]: array([[2005.0, 179.0, '150000', ..., 1, 1, 4],
 [1997.0, 60.0, '150000', ..., 38, 7, 4],
 [2003.0, 170.0, '150000', ..., 2, 7, 1],
 ...,
 [2009.0, 174.0, '125000', ..., 25, 7, 7],
 [2000.0, 136.0, '150000', ..., 20, 7, 1],
 [2013.0, 170.0, '40000', ..., 1, 7, 8]], dtype=object)

```
In [27]: regressor.predict([[2012.0, 179.0, '150000', 12.0, 0, 0, 30, 1, 1, 4]])
```

Out[27]: array([18518.37919135])